

Word Embeddings

Rohit Amarnath

December 5, 2023

1 Continous Bag of Words

The CBOW model takes a window of surrounding words as input and tries to predict the target word in the center of the window.

Let V be the vocabulary, n be the embedding size, $E = \begin{bmatrix} | & & | \\ e_1 & \dots & e_{|V|} \\ | & & | \end{bmatrix} \in \mathbb{R}^{n \times |V|}$ be the input word matrix and $U = \begin{bmatrix} - & u_1 & - \\ & \vdots & \\ - & u_{|V|} & - \end{bmatrix} \in \mathbb{R}^{|V| \times n}$ be the output word matrix. The input word matrix can alternatively be interpreted as the word embedding layer. Each column in this matrix is a word embedding.

1.1 Methodology

1. The input is a $2m$ size list of words $(x_{c-m}, \dots, x_{c-1}, x_{c+1}, \dots, x_{c+m})$. Each word is a number, and we will one hot encode these vectors before feeding them into our model. The output is $x_c = k$, and it is also one hot encoded.
2. We get the embedded word vectors

$$(v_{c-m}, \dots, v_{c-1}, v_{c+1}, \dots, v_{c+m}) = (Ex_{c-m}, \dots, Ex_{c-1}, Ex_{c+1}, \dots, Ex_{c+m})$$

Then, average the embedded vectors.

$$\bar{v} = \frac{v_{c-m} + \dots + v_{c-1} + v_{c+1} + \dots + v_{c+m}}{2m}$$

3. Generate a score vector $z = U\bar{v}$ and turn the scores into probabilities with the softmax function. We have $\hat{y} = \text{softmax}(z)$.
4. Cross-entropy loss is used to train the model.

$$H(\hat{y}, y) = - \sum_{j=1}^{|V|} y_j \log(\hat{y}_j)$$

Since $x_c = k$ as defined in the first step, the loss is $H(\hat{y}, y) = -y_c \log(\hat{y}_c)$. Thus, with gradient descent, we optimize

$$\begin{aligned} H(\hat{y}, y) &= -\log \hat{y}_k \\ &= -\log \frac{\exp(u_k^T \hat{v})}{\sum_{j=1}^{|V|} \exp(u_j^T \hat{v})} \\ &= -u_k^T \hat{v} + \log \sum_{j=1}^{|V|} \exp(u_j^T \hat{v}) \end{aligned}$$

2 Skip-grams

The skip-gram is the inverse of the CBOW model. In the skip-gram, the input is a word, and the model tries to predict the context window surrounding the word.

As before, let V be the vocabulary, n be the embedding size, $E = \begin{bmatrix} | & & | \\ e_1 & \dots & e_{|V|} \\ | & & | \end{bmatrix} \in \mathbb{R}^{n \times |V|}$ be the input word matrix and $U = \begin{bmatrix} - & u_1 & - \\ & \vdots & \\ - & u_{|V|} & - \end{bmatrix} \in \mathbb{R}^{|V| \times n}$ be the output word matrix.

2.1 Methodology

There is an abuse of notation with the indexing in this section. When we index by i , we actually index by the word index. Let $x_i = k$, which means x_i is the k -th word in the vocabulary. The notation used in this section allows $u_i = u_{x_i}$ and $e_i = e_{x_i}$.

The skip-gram model learns the probability $P(x_{c-m}, \dots, x_{c-1}, x_{c+1}, \dots, x_{c+m} \mid x_c)$. We invoke the naive Bayes assumption and assume conditional independence between context words. That is,

$$P(x_{c-m}, \dots, x_{c-1}, x_{c+1}, \dots, x_{c+m} \mid x_c) = \prod_{i=-m, i \neq 0}^m P(x_{c+i} \mid x_c)$$

We want to minimize the negative of the above. The skip-gram model can be thought of as training a hidden representation (the word embedding) by training a softmax regression on $2m - 1$ outputs for one input.

We now go over the model.

1. The input is a word. As before, the word is a number, but we will one hot encode this before feeding it into the model. Let the input be x_c and the output be $x_{c-m}, \dots, x_{c-1}, x_{c+1}, \dots, x_{c+m}$.
2. We retrieve the embedded word vector $v_c = Ex_c$. Now, for each word in the context, we calculate the softmax probability. That is, for the output word x_{c-m} , we calculate

$$\begin{aligned} P(x_{c-m} \mid x_c) &= P(u_{c-m} \mid v_c) \\ &= \frac{\exp(u_{c-m}^T v_c)}{\sum_{i=1}^{|V|} \exp(u_i^T v_c)} \end{aligned}$$

We now have a set of probabilities for every word in the context.

3. With gradient descent, we minimize the negative log of $P(x_{c-m}, \dots, x_{c-1}, x_{c+1}, \dots, x_{c+m} \mid x_c)$.

$$\begin{aligned}
-\log P(x_{c-m}, \dots, x_{c-1}, x_{c+1}, \dots, x_{c+m} \mid x_c) &= -\log \prod_{i=-m, i \neq 0}^m P(x_{c+i} \mid x_c) \\
&= -\log \prod_{i=-m, i \neq 0}^m P(u_{c+i} \mid v_c) \\
&= -\sum_{i=-m, i \neq 0}^m \log P(u_{c+i} \mid v_c) \\
&= -\sum_{i=-m, i \neq 0}^m \log \frac{\exp(u_{c+i}^T v_c)}{\sum_{i=1}^{|V|} \exp(u_i^T v_c)} \\
&= -\sum_{i=-m, i \neq 0}^m u_{c+i}^T v_c + 2m \log \sum_{i=1}^{|V|} \exp(u_i^T v_c)
\end{aligned}$$

We can also think of this as finding the cross entropy loss with respect to $2m - 1$ output vectors, showing the correspondence between cross entropy loss and the maximum likelihood.

3 Negative Sampling

Skip-grams work well, but they require a $O(|V|)$ update per word and context pair. To mitigate this, we reframe the prediction problem from maximizing likelihoods of pairs of context and a word to determining whether a pair of words is a pair of context and a word. Observe that this new paradigm has changed the existing $|V|$ classes classification task to a binary classification task (whether a pair of words is *valid* or not). Negative sampling is a special case of noise contrastive estimation and is a contrastive representation learning technique. We present it now.

Denote $(w, c) \in \mathcal{D}$ be a pair of word and context. Let $P(D = 1 \mid w, c)$ be the probability that pair (w, c) came from the corpus.

$$P(D = 1 \mid w, c) = \frac{1}{1 + \exp(-u_w^T v_c)}$$

Similarly, $P(D = 0 \mid w, c)$ is the probability that pair (w, c) did not come from the corpus.

Let $\tilde{\mathcal{D}}$ be the set of pairs of word and context that are erroneous and did not come from the corpus. With these probabilities, we have the maximum likelihood to be

$$\begin{aligned}
\theta_{\text{MLE}} &= \arg\max_{\theta} \prod_{(w,c) \in \mathcal{D}} P(D = 1 \mid w, c) \prod_{(w,c) \in \tilde{\mathcal{D}}} P(D = 0 \mid w, c) \\
&= \arg\max_{\theta} \sum_{(w,c) \in \mathcal{D}} \log P(D = 1 \mid w, c) + \sum_{(w,c) \in \tilde{\mathcal{D}}} \log(1 - P(D = 1 \mid w, c)) \\
&= \arg\max_{\theta} \sum_{(w,c) \in \mathcal{D}} \log \frac{1}{1 + \exp(-u_w^T v_c)} + \sum_{(w,c) \in \tilde{\mathcal{D}}} \log \left(1 - \frac{1}{1 + \exp(-u_w^T v_c)} \right) \\
&= \arg\max_{\theta} \sum_{(w,c) \in \mathcal{D}} \log \frac{1}{1 + \exp(-u_w^T v_c)} + \sum_{(w,c) \in \tilde{\mathcal{D}}} \log \frac{1}{1 + \exp(u_w^T v_c)}
\end{aligned}$$

The objective function ends up being

$$J = -\log \frac{1}{1 + \exp(-u_w^T v_c)} - \sum_{i=1}^K \log \frac{1}{1 + \exp(u_i^T v_c)}$$

where we take k samples from $\tilde{\mathcal{D}}$ and u_1, \dots, u_k are the output embeddings of the *negative samples* and v_c is the input word. Note that $(w, c) \in \mathcal{D}$ and $(i, c) \in \tilde{\mathcal{D}}$.

4 GloVe: Global Vectors

Denote X to be the co-occurrence matrix of vocabulary V . X is a $|V| \times |V|$ matrix, and X_{ij} denotes the count of word i being within the context of word j and vice versa (note that $X_{ij} = X_{ji}$). The GloVe algorithm has two $K \times N$ matrices W and \tilde{W} which factorize the co-occurrence matrix X into a K dimensional subspace.

The GloVe algorithm solves the following prediction problem:

$$W_i^T \tilde{W}_j + b_i + \tilde{b}_j = \log X_{ij}$$

where W_i^T is the i -th column of W , \tilde{W}_j is the j -th column of \tilde{W} , and b_i and \tilde{b}_j are bias terms for words i and j , respectively.

The loss function is

$$J = \sum_{i=1}^{|V|} \sum_{j=1}^{|V|} f(X_{ij}) (W_i^T \tilde{W}_j + b_i + \tilde{b}_j - \log X_{ij})^2$$

The GloVe paper has a more specific derivation of the loss function, which we will not present here. After training, we compute $\frac{W + \tilde{W}}{2}$ to get the word embeddings.